

NA 568 - Winter 2022

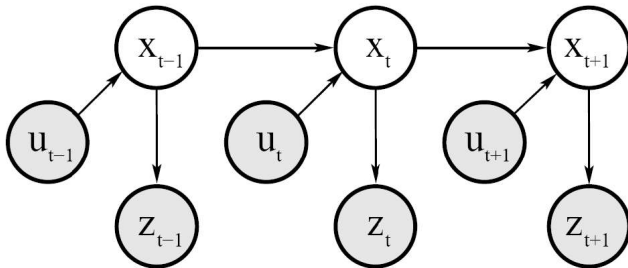
Nonlinear Kalman Filtering

Maani Ghaffari

January 13, 2021



Dynamic Bayesian Network for Controls, States, and Sensations



State Estimation

- Estimate the state x of a system given observations z and controls u
- **Goal:**

$$p(x \mid z, u)$$

Courtesy: C. Stachniss

Prediction and Correction Step

- Bayes filter can be written as a two step process

- **Prediction step**

$$\overline{bel}(x_t) = \int p(x_t \mid u_t, x_{t-1}) bel(x_{t-1}) dx_{t-1}$$

- **Correction step**

$$bel(x_t) = \eta p(z_t \mid x_t) \overline{bel}(x_t)$$

Courtesy: C. Stachniss

Motion and Observation Model

□ Prediction step

$$\overline{bel}(x_t) = \int \underbrace{p(x_t \mid u_t, x_{t-1})}_{\text{motion model}} bel(x_{t-1}) dx_{t-1}$$

□ Correction step

$$bel(x_t) = \eta \underbrace{p(z_t \mid x_t)}_{\text{sensor or observation model}} \overline{bel}(x_t)$$

Courtesy: C. Stachniss

▶ Given:

- ▶ Stream of observations $z_{1:t}$ and action data $u_{1:t}$
- ▶ Sensor/measurement model $p(z_t|x_t)$
- ▶ Action/motion/transition model $p(x_t|x_{t-1}, u_t)$

▶ Wanted:

- ▶ The state x_t of dynamical system
- ▶ The posterior of state is called belief $bel(x_t) = p(x_t|z_{1:t}, u_{1:t})$

Algorithm 1 Bayes-filter

Require: Belief $bel(x_{t-1}) = p(x_{t-1}|z_{1:t-1}, u_{1:t-1})$, action u_t , measurement z_t ;

1: **for** all state variables **do**

2: $\overline{bel}(x_t) = \int p(x_t|x_{t-1}, u_t)bel(x_{t-1})dx_{t-1}$ // Predict using action/control input u_t

3: $bel(x_t) = \eta p(z_t|x_t)\overline{bel}(x_t)$ // Update using perceptual data z_t

4: **return** $bel(x_t)$

Nonlinear dynamic system with additive noise.

- ▶ Nonlinear process model:

$$x_k = f(u_k, x_{k-1}) + w_k$$

- ▶ Nonlinear measurement model:

$$z_k = h(x_k) + v_k$$

Nonlinear dynamic system with multiplicative noise.

- ▶ Nonlinear process model:

$$x_k = f(u_k, x_{k-1}, w_k)$$

- ▶ Nonlinear measurement model:

$$z_k = h(x_k, v_k)$$

Q. How to map belief (a probability distribution) through a nonlinear function?

Q. How to map belief (a probability distribution) through a nonlinear function?

Key ideas:

- ▶ Linearization via Taylor expansion

Q. How to map belief (a probability distribution) through a nonlinear function?

Key ideas:

- ▶ Linearization via Taylor expansion
- ▶ Unscented Transform (deterministic sampling)

Q. How to map belief (a probability distribution) through a nonlinear function?

Key ideas:

- ▶ Linearization via Taylor expansion
- ▶ Unscented Transform (deterministic sampling)
- ▶ Monte-Carlo methods (random sampling)

Q. How to map belief (a probability distribution) through a nonlinear function?

Key ideas:

- ▶ Linearization via Taylor expansion
→ Extended Kalman Filter (EKF)
- ▶ Unscented Transform (deterministic sampling)
→ Unscented Kalman Filter (UKF)
- ▶ Monte-Carlo methods (random sampling)
→ Sequential Monte-Carlo methods (Particle Filters)

Linearization via Taylor Expansion

Linearization of a function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ around point a is

$$\begin{aligned} f(x) &\approx f(a) + \left. \frac{\partial f}{\partial x} \right|_{x=a} (x - a) \\ &= \left(f(a) - \left. \frac{\partial f}{\partial x} \right|_{x=a} a \right) + \left. \frac{\partial f}{\partial x} \right|_{x=a} x \\ &:= x_0 + Fx \end{aligned}$$

Affine! We know how to propagate a Gaussian through an affine map.

Recall: Affine Transformation of a Multivariate Gaussian

Suppose $x \sim \mathcal{N}(\mu, \Sigma)$ and $y = Ax + b$.

Then $y \sim \mathcal{N}(A\mu + b, A\Sigma A^T)$.

$$F_k = \left. \frac{\partial f}{\partial x} \right|_{x=\mu_{k-1}}, W_k = \left. \frac{\partial f}{\partial w} \right|_{x=\mu_{k-1}}, H_k = \left. \frac{\partial h}{\partial x} \right|_{x=\mu_k^-}, V_k = \left. \frac{\partial h}{\partial v} \right|_{x=\mu_k^-}$$

Algorithm 2 Extended-Kalman-filter

Require: belief mean μ_{k-1} , belief covariance Σ_{k-1} , action u_k , measurement z_k ;

- 1: $\mu_k^- \leftarrow f(u_k, \mu_{k-1})$ ▷ predicted mean
 - 2: $\Sigma_k^- \leftarrow F_k \Sigma_{k-1} F_k^\top + W_k Q_k W_k^\top$ ▷ predicted covariance
 - 3: $\nu_k \leftarrow z_k - h(\mu_k^-)$ ▷ innovation
 - 4: $S_k \leftarrow H_k \Sigma_k^- H_k^\top + V_k R_k V_k^\top$ ▷ innovation covariance
 - 5: $K_k \leftarrow \Sigma_k^- H_k^\top S_k^{-1}$ ▷ filter gain
 - 6: $\mu_k \leftarrow \mu_k^- + K_k \nu_k$ ▷ corrected mean
 - 7: $\Sigma_k \leftarrow (I - K_k H_k) \Sigma_k^-$ ▷ corrected covariance
 - 8: $// \Sigma_k \leftarrow (I - K_k H_k) \Sigma_k^- (I - K_k H_k)^\top + K_k R_k K_k^\top$ ▷ numerically stable form
 - 9: **return** μ_k, Σ_k
-

Example: EKF Target Tracking

A target is moving in a 2D plane. The ownship position is known and fixed at the origin. We have access to relative noisy range and bearing measurements of the target position at any time step.

$$x_k = f(u_k, x_{k-1}) + w_k = x_{k-1} + w_k$$


$$z_k = h(x_k) + v_k = \begin{bmatrix} \sqrt{x_k^1{}^2 + x_k^2{}^2} \\ \text{atan2}(x_k^1, x_k^2) \end{bmatrix} + v_k$$

$$F_k = I_2, G_k = 0_2, Q_k = 0.001 I_2, R_k = \text{diag}(0.05^2, 0.01^2)$$

$$H_k = \begin{bmatrix} \frac{x_k^1}{\sqrt{x_k^1{}^2 + x_k^2{}^2}} & \frac{x_k^2}{\sqrt{x_k^1{}^2 + x_k^2{}^2}} \\ \frac{x_k^2}{x_k^1{}^2 + x_k^2{}^2} & \frac{-x_k^1}{x_k^1{}^2 + x_k^2{}^2} \end{bmatrix}$$

Example: EKF Target Tracking

See `ekf_single_target.m` for code.



- ▶ Highly efficient; polynomial time in measurement dimensionality n_z and state dimensionality n_x .
- ▶ Not optimal.
- ▶ Can diverge if nonlinearities are large.
- ▶ Can work well in practice for many problems despite violating all the underlying assumptions.

Unscented Transform Overview

- ▶ Compute a set of sigma points (samples) \mathcal{X} ;
- ▶ Each sigma point has a weight w ;
- ▶ Transform the point through the nonlinear function $g(x)$;
- ▶ Compute a Gaussian distribution from weighted points using weighted sample mean and covariance;

The first sigma point is the mean.

$$x_0 = \mu$$

$$x_i = \mu + \ell'_i \quad i = 1, \dots, n$$

$$x_i = \mu - \ell'_{i-n} \quad i = n+1, \dots, 2n$$

ℓ'_i is the i -th column of L' where $L' = \sqrt{(n + \kappa)L}$ and $\Sigma = LL^T$ can be computed using Cholesky decomposition. n is the dimension of the state and κ is a user-definable parameter.

Let $g : \mathbb{R}^n \rightarrow \mathbb{R}^m$. Compute a Gaussian from weighted and transformed points

$$\mu' = \sum_{i=0}^{2n} w_i g(x_i)$$

$$\Sigma' = \sum_{i=0}^{2n} w_i (g(x_i) - \mu')(g(x_i) - \mu')^\top$$

$$w_0 = \frac{\kappa}{n + \kappa}$$

$$w_i = \frac{1}{2(n + \kappa)} \quad i = 1, \dots, 2n$$

The user-defined parameter, κ , can be tuned to adjust the weight for a particular transformation. For instance $\kappa = 2$ (see *State Estimation for Robotics*, Timothy D. Barfoot, 2018, Ch. 4.2.7.).

Remark

If the noise is additive, we simply add the noise covariance to the propagated sample covariance. If the noise is multiplicative, we augment the state with noise by adding zeros of appropriate dimension to the mean and constructing a block-diagonal covariance matrix of the state and noise covariances. Note that in the latter case the dimension of the augmented state is increased to the sum of the dimensions of the state and noise vectors; hence, more samples need to be drawn. See State Estimation for Robotics, Timothy D. Barfoot, 2018, Ch. 4.2.9.

Example: Unscented Transform



See `unscented_transform_example.m` for code.

Transform a Gaussian distribution from polar to Cartesian coordinates.

$$\begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} r \cos(\theta) \\ r \sin(\theta) \end{bmatrix}, \quad \begin{bmatrix} r \\ \theta \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} 1.5 \\ \pi/6 \end{bmatrix}, \begin{bmatrix} 0.3^2 & -0.14^2 \\ -0.14^2 & 0.35^2 \end{bmatrix}\right)$$



Example: Unscented Transform

$$\begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} r \cos(\theta) \\ r \sin(\theta) \end{bmatrix}, \quad \begin{bmatrix} r \\ \theta \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} 1.5 \\ \pi/6 \end{bmatrix}, \begin{bmatrix} 0.3^2 & -0.14^2 \\ -0.14^2 & 0.35^2 \end{bmatrix}\right)$$



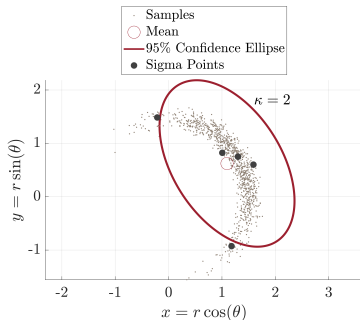

Example: Unscented Transform

$$\begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} r \cos(\theta) \\ r \sin(\theta) \end{bmatrix}, \quad \begin{bmatrix} r \\ \theta \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} 1.5 \\ \pi/6 \end{bmatrix}, \begin{bmatrix} 0.3^2 & -0.14^2 \\ -0.14^2 & 0.35^2 \end{bmatrix}\right)$$



Example: Unscented Transform

$$\begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} r \cos(\theta) \\ r \sin(\theta) \end{bmatrix}, \quad \begin{bmatrix} r \\ \theta \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} 1.5 \\ \pi/6 \end{bmatrix}, \begin{bmatrix} 0.1^2 & -0.09^2 \\ -0.09^2 & 0.6^2 \end{bmatrix}\right)$$



Algorithm 3 Unscented-Kalman-filter

Require: belief mean μ_{k-1} , belief covariance Σ_{k-1} , action u_k , measurement z_k ;

- 1: $\mathcal{X}_{k-1} \leftarrow$ compute the set of $2n + 1$ sigma points using μ_{k-1} and Σ_{k-1}
 - 2: $w^- \leftarrow$ compute the set of $2n + 1$ weights
 - 3: $\mu_k^- = \sum_{i=0}^{2n} w_i^- f(u_k, x_{k-1,i})$ ▷ predicted mean
 - 4: $\Sigma_k^- \leftarrow \sum_{i=0}^{2n} w_i^- (f(u_k, x_{k-1,i}) - \mu_k^-)(f(u_k, x_{k-1,i}) - \mu_k^-)^\top + Q_k$ ▷ predicted covariance
 - 5: $\mathcal{X}_k^- \leftarrow$ compute the set of $2n + 1$ sigma points using μ_k^- and Σ_k^-
 - 6: $w \leftarrow$ compute the set of $2n + 1$ weights
 - 7: $z_k^- = \sum_{i=0}^{2n} w_i h(x_{k,i}^-)$ ▷ predicted measurement
 - 8: $\nu_k \leftarrow z_k - z_k^-$ ▷ innovation
 - 9: $S_k \leftarrow \sum_{i=0}^{2n} w_i (h(x_{k,i}^-) - z_k^-)(h(x_{k,i}^-) - z_k^-)^\top + R_k$ ▷ innovation covariance
 - 10: $\Sigma_k^{xz} \leftarrow \sum_{i=0}^{2n} w_k^{[i]} (x_{k,i}^- - \mu_k^-)(h(x_{k,i}^-) - z_k^-)^\top$ ▷ state and measurement cross covariance
 - 11: $K_k \leftarrow \Sigma_k^{xz} S_k^{-1}$ ▷ filter gain
 - 12: $\mu_k \leftarrow \mu_k^- + K_k \nu_k$ ▷ corrected mean
 - 13: $\Sigma_k \leftarrow \Sigma_k^- - K_k S_k K_k^\top$ ▷ corrected covariance
 - 14: **return** μ_k, Σ_k
-

Example: UKF Target Tracking

A target is moving in a 2D plane. The ownship position is known and fixed at the origin. We have access to relative noisy range and bearing measurements of the target position at any time step.


$$x_k = f(u_k, x_{k-1}) + w_k = x_{k-1} + w_k$$

$$z_k = h(x_k) + v_k = \begin{bmatrix} \sqrt{x_k^1{}^2 + x_k^2{}^2} \\ \text{atan2}(x_k^1, x_k^2) \end{bmatrix} + v_k$$

$$F_k = I_2, G_k = 0_2, Q_k = 0.001 I_2, R_k = \text{diag}(0.05^2, 0.01^2)$$

Example: UKF Target Tracking

See `ukf_single_target.m` for code.



- ▶ Highly efficient: same complexity as EKF, with a constant factor slower in typical practical applications;
- ▶ Better linearization than EKF;
- ▶ Derivative-free: no Jacobians needed.
- ▶ Not optimal.

- ▶ Same results as EKF for linear models;
- ▶ Better approximation than EKF for non-linear models;
- ▶ Differences often “somewhat small”;
- ▶ No Jacobians needed for the UKF;
- ▶ Same complexity class;
- ▶ Slightly slower than the EKF

- ▶ Probabilistic Robotics: Ch. 3
- ▶ State Estimation for Robotics: Ch. 3 and 4